

The Imagination of Nature:

Mythology as the Cultural Recording Layer of Perceptual Coupling

Antoine Bellemare-Pepin¹ and Mar Estarellas²

¹Department of Psychology, Université de Montréal

²Department of Psychiatry, McGill University

June 3, 2026

Abstract

Mythography has long treated the symbolic content of myth as a function of mind and culture, with the biome supplying raw material but not form. A counter-tradition from Schelling through Bateson and the ecological anthropology of Descola and Ingold treats imagination instead as a coupling between perceiver and environment, on which mythologies should carry a recoverable trace of their biomes. We test this with Berezkin's catalogue of folk-motifs from 958 traditions, geocoded to the WWF terrestrial biomes and aligned via sentence-pooled SigLIP-2 to iNaturalist species photographs and Places365 landscape scenes after an LLM pipeline strips species, place, ethnonym, and biome-word vocabulary from the motif text. The alignment survives on both image corpora, replicates across four vision–language models, and decays monotonically with motif breadth, strongest in biome-specific motifs and fading toward zero in universals. The pattern is consistent with perceptual coupling at the lexical-thematic level. The residue mythology carries of its biome is not a catalogue of names but the schemas of attention, activity, and form that the biome made available to imagination. On this reading, mythology functions as the cultural recording layer of an ongoing engagement between perceiver and environment, a substrate that survives even when species, place, and biome words are stripped from the text because the perceptual encounter shaped the very activities and articulations the narrative is built from. The breadth gradient is the load-bearing signature. Biome-specific motifs stabilised in encounter with a particular ecology and carry that encounter forward into their narrative substance, while universal motifs propagated only by detaching from any single biome and stabilising on cross-cultural structures, so the coupling fades in them precisely as the lineage from Schelling and Bateson through contemporary 4E cognitive science would predict. Mythology, on the reading the data support, is one of the long instruments humans have used to keep the coupling between cognition and environment articulate across generations.

Significance Statement

Mythologies of cultures inhabiting different biomes carry a recoverable visual signature of those biomes, even after every species name, place name, and biome word is stripped from the text. The signature is strongest in biome-specific myths and attenuates progressively as motifs propagate across more biomes. The pattern is consistent with the hypothesis that the ecological structure of a biome contributes to the stabilisation of its mythological motifs, and reopens, in empirically tractable form, a two-century-old debate over whether imagination is a private projection of cultural meaning onto a passive world or a coupling phenomenon between perceiver and environment.

Introduction

Modern thought has long treated nature as a passive backdrop onto which the human mind projects meaning. Mythology, on this view, is what humans *do to* a world, a projection of cultural categories,

anxieties, and social structures onto matter. The biome a culture inhabits supplies raw material but not, on this reading, the *form* of its symbols. We call this dominant view the *projection account*. Running parallel to this default, however, a different tradition has insisted on a different picture. Schelling’s distinction between *natura naturata* (nature as finished product) and *natura naturans* (nature as productive activity) [Schelling, 1797, 1799] was recovered in the twentieth century by Bateson’s claim that “mind is wherever the conditions for a mental process obtain” across coupled systems, where differences make differences and pattern propagates between brain, body, environment, and culture [Bateson, 1972, 1979]. Depth psychology made a related move. Jung’s mature view of the archetypes resisted easy translation into innate mental contents stored in individual brains and instead located them in the joint between psyche and nature [Jung, 1959]; Hillman radicalised this into archetypal psychology and the recovery of *anima mundi* [Hillman, 1975, 1981]; Henri Corbin gave it its most precise philosophical name in the *mundus imaginalis*, the imaginal world, distinct from the modern collapse of “imaginary” into “merely subjective” [Corbin, 1972, Barfield, 1957, Abram, 1996, Ingold, 2000]. The thread connecting these authors is that imagination is a perceptual organ rather than a productive one, an organ that discloses content present in the encounter between an attentive perceiver and an articulate world. We call this opposing view the *coupling account*, since on this reading perceiver and environment co-produce symbolic content together.

Since the 1960s the cognitive sciences have been producing the empirical scaffolding for exactly this coupling view. Gibson’s ecological psychology replaced the dominant representational model of perception with one of *direct pickup*, in which the environment is information-rich at the level of the optic array and the organism tunes itself to a richly structured world that already contains the relevant information [Gibson, 1979]. The 4E framework (embodied, embedded, extended, and enactive cognition) generalised this beyond perception [Varela et al., 1991, Clark and Chalmers, 1998, Noë, 2004]. Predictive processing supplied the neural substrate, with the brain treated as a hierarchical prediction machine that locks onto interpretations its priors privilege, given the actual statistical structure of what the world presents [Friston, 2010, Clark, 2016, Hohwy, 2013]. On this account pareidolia is not a mistake of perception but the system working as designed. Niche construction theory closed the loop [Odling-Smee et al., 2003], observing that humans produce symbolic environments (names, stories, transmitted practices) and that these environments shape the cognition of subsequent generations.

The two accounts make different empirical predictions about mythology. The projection account treats the symbolic content of myth as largely a product of mind and culture, and therefore predicts that once the text of a mythology has been stripped of its overt environmental and cultural anchors (species names, place names, ethnonyms, biome words, climate phrases, tradition labels), the residual narrative should carry no detectable visual signature of the biome the tradition inhabits. On this reading, what remains of mythology after anonymisation is the cross-cultural grammar of plot and relation, and that grammar should be indifferent to which biome a tradition emerged in. The coupling account treats imagination instead as a perceptual organ jointly tuned by mind and environment, and therefore predicts that the structural character of a biome, its light, its dominant forms, its taxa, the affordances it presents to a perceiver, should leave a recoverable residue in the surviving text, because that residue was part of what stabilised the mythology in the first place. A secondary prediction of the coupling account, which the projection account does not share, is that this residue should be strongest in motifs whose distribution remained close to a single biome and attenuate as motifs propagate across many biomes, since universally-shared motifs are precisely those whose stabilising encounters were not tied to one environment.

Two recent developments make this question newly testable at scale. The first is the cross-cultural folk-motif catalogue assembled by Yu. E. Berezkin, comprising 2,564 motifs across 958 traditions with Russian-language abstracts and geocoded centroids [Berezkin, 2024], which we have joined to the WWF Ecoregions classification of 14 terrestrial biomes [Olson et al., 2001]. The second is the development of frontier vision–language models such as SigLIP-2 [Tschannen et al., 2025], which embed text and image in a shared metric space learned from billions of image–text pairs sampled from the internet. The combination lets us ask whether, after stripping every species name, every place name, every biome word, and every tradition label, the embedding of a tradition’s mythology preferentially aligns with images of its

own biome, and whether that alignment is concentrated in biome-specific motifs as the coupling account predicts.

We report below that such a residue is detectable. After stripping every overt environmental and cultural anchor from the text, the embedding of a tradition’s mythology preferentially aligns with images of its own biome on two independent image corpora and across four vision–language model architectures, in both Russian and English. The alignment is graded with motif breadth as the coupling account predicts, strongest in motifs that did not propagate beyond a single biome and attenuating in motifs shared across many. We take this as evidence consistent with the coupling account and not naturally predicted by the projection account, and read mythology as the surviving textual residue of an ongoing perceptual coupling between cognition and environment, the recoverable trace in transmitted narrative of a living engagement that the narrative itself does not exhaust.

Results

Biome-specific mythology aligns with biome imagery




Figure 1. Cross-cultural sample of mythology over the world’s biome geography. World map of the 958 Berezkin traditions plotted as dots over the underlying WWF Ecoregions polygons, both coloured by the same biome palette: continents in semi-transparent biome fills (the Saharan and Australian deserts in tan, the Amazonian and Indo-Malay rainforests in dark green, the Boreal belt across northern Eurasia and Canada in muted blue, Mediterranean basins in red, etc.), tradition centroids as opaque dots in the same palette. The overlay shows how densely each biome’s actual extent is sampled by Berezkin’s corpus. Eurasian and North American traditions are relatively dense; some biomes (Mangroves, Flooded Grasslands and Savannas) are sparsely covered.

The Berezkin corpus distributes 958 traditions across the 14 WWF biomes (Figure 1). After the full anonymisation pipeline described in *Materials and Methods*, the LLM-mediated translation and double-pass refinement that replaces every species with a class word, every place and ethnonym with a placeholder, and every tradition-specific character name and cultural-religious title with a neutral generic, the residualised alignment on the full LLM-clean corpus (2,158 motifs, sentence-pooled SigLIP-2) reaches within-iconic-taxon stratified $\mu\Delta = +0.406 \times 10^{-3}$ across the 14 biomes on iNaturalist (Figure 2, Panel A; 8/14 biomes significant at $p < 0.05$). The strongest positive cells under stratification lie in Montane Grasslands & Shrublands ($\Delta^{\text{strat}} = +1.08 \times 10^{-3}$, $p < .001$), Tropical & Subtropical Moist Broadleaf Forests ($+0.93 \times 10^{-3}$, $p < .001$), Tropical & Subtropical Grasslands, Savannas & Shrublands ($+0.76 \times 10^{-3}$, $p < .001$), Boreal Forests/Taiga ($+0.65 \times 10^{-3}$, $p < .001$), Deserts & Xeric Shrublands ($+0.63 \times 10^{-3}$, $p < .001$), Tropical & Subtropical Dry Broadleaf Forests ($+0.44 \times 10^{-3}$, $p < .01$), Tundra ($+0.36 \times 10^{-3}$, $p < .01$), and Temperate Grasslands, Savannas & Shrublands ($+0.33 \times 10^{-3}$,




Figure 2. The headline coupling signature, replicated on two independent image corpora. Residualised Δ per WWF biome under sentence-pooled SigLIP-2 on the full LLM-clean Berezkin abstracts (2,158 motifs). *Panel A:* iNaturalist species photographs ($n = 47,478$ images), with Δ stratified within each iconic-taxon stratum and uniformly averaged across taxa to remove biome-by-iconic-taxon co-occurrence as a possible explanation; stratified $\mu\Delta = +0.406 \times 10^{-3}$, 8/14 biomes significant at $p < .05$. *Panel B:* Places365 strict-landscape scenes ($n = 1,675$ images, 11 biomes); because Places365 has no iconic-taxon labels (it is scene-only imagery), we report the marginal residualised Δ ; marginal $\mu\Delta = +0.423 \times 10^{-3}$, 6/11 biomes significant at $p < .05$. Biomes are ordered top-to-bottom by descending iNat stratified Δ ; biomes Places365 does not cover are blanked. Each panel uses its own dynamic range. Gold bold asterisks denote significance after Benjamini-Hochberg FDR correction ($q < .05$); grey asterisks denote nominal $p < .05$. The two corpora agree on the biomes that drive the effect (Tropical Moist Broadleaf, Deserts, Boreal/Taiga, Montane Grasslands, Tropical Grasslands), an independent landscape-channel replication of a signal first recovered on species-channel photographs.

$p < .01$). The remaining biomes are positive or near zero with the exception of Mediterranean Forests and Temperate Conifer Forests, both weakly negative under stratification.

The biomes that survive the within-iconic-taxon control therefore carry alignment that is not attributable to biome-by-iconic-taxon co-occurrence on the image side; the cells where stratification narrows the marginal effect precisely identify where the marginal signal was riding on taxonomic composition rather than on biome-character imagery. Mediterranean Forests is the cleanest example: its marginal $\Delta = +0.51 \times 10^{-3}$ ($p < .001$) is carried by a single strongly positive Plantae cell ($+1.38 \times 10^{-3}$ under the per-taxon decomposition, Figure 5), because Mediterranean iNaturalist imagery is dominated by plants (vineyards, olive groves, shrubland) and the motifs assigned to Mediterranean biomes lean similarly plant-coded. Once we hold iconic-taxon fixed on the image side, the Plantae contribution is uniformly averaged against the negative Aves (-0.81), Animalia (-1.12), Reptilia (-0.39), and Amphibia (-0.91) cells, and the stratified Δ collapses to -0.09×10^{-3} (panel A bottom). The reverse case is also visible. Boreal Forests/Taiga has a near-zero marginal Δ ($+0.03 \times 10^{-3}$) because its dominant Plantae stratum (conifer imagery) carries no distinctive boreal mythological content; once stratification gives Mammalia (bears, wolves, $+1.77 \times 10^{-3}$) and Actinopterygii ($+1.39$) equal weight, the boreal Δ emerges at $+0.65 \times 10^{-3}$ ($p < .001$).

The same pattern replicates on an independent image corpus. The Places365 strict-landscape comparison set, 1,675 scene photographs across 11 biomes drawn from Places365 scene categories that unambiguously match a single WWF biome and then SigLIP-filtered against humans, vehicles, buildings, and close-up animal subjects, gives a marginal $\mu\Delta = +0.423 \times 10^{-3}$ across 11 biomes, with 6




Figure 3. Per-biome Δ projected onto the world map. Geographic companion to Figure 2, on the within-iconic-taxon stratified Δ under sentence-pooled SigLIP-2. *Panel A*: full LLM-clean corpus (2,158 motifs), matching the headline statistic of Figure 2, Panel A. *Panel B*: Spec A subset (motifs touching ≤ 3 biomes), the breadth-restricted view that the coupling reading predicts should show the cleanest geography (cf. Figure 4). Colour intensity scales with stratified Δ further weighted by confidence ($1 - p$); the YlOrRd colormap is shared between panels so cell intensities are directly comparable. Light grey marks sub-confidence biomes (fewer than 10 traditions, 50 motifs, or 50 photographs); muted slate-blue marks the rare biomes whose stratified Δ falls below zero (in panel A, Mediterranean Forests and Temperate Conifer Forests). Black dots overlay the 958 tradition centroids. The headline coupling geography appears in both panels (Tropical Moist Broadleaf rainforests of the Amazon and Indo-Malay, Tropical Grasslands of sub-Saharan Africa, Boreal Forests/Taiga of northern Eurasia and Canada, Deserts of North Africa and central Australia, Montane Grasslands of the high-altitude tropics) and is substantially more vivid in the Spec A panel where breadth dilution is removed.

biomes significant at $p < 0.05$ (Figure 2, Panel B). The strongest cells fall in Tropical & Subtropical Moist Broadleaf Forests ($\Delta = +1.67 \times 10^{-3}$, $p < .001$), Mediterranean Forests, Woodlands & Scrub ($+1.25 \times 10^{-3}$, $p < .001$), Boreal Forests/Taiga ($+1.06 \times 10^{-3}$, $p < .001$), Deserts & Xeric Shrublands ($+0.94 \times 10^{-3}$, $p < .001$), and Montane Grasslands & Shrublands ($+0.57 \times 10^{-3}$, $p < .001$). Crucially, the two corpora cover different visual content: iNaturalist samples close to species-level (foliage, fauna, fungi in their habitat), Places365 samples wide-shot terrain with no species or human subjects. The fact that anonymised mythology aligns above permutation chance with both channels of biome imagery, species-level and scene-level, means the alignment is not an artefact of one corpus’s curatorial peculiarities. Projected onto a world map (Figure 3), the warmest cells concentrate in the Amazonian and Indo-Malay Tropical Moist Broadleaf rainforests, the Saharan and Australian deserts, the Mediterranean basins of southern Europe and the western Maghreb, the boreal belt across northern Eurasia and Canada, and the montane grassland systems of the high-altitude tropics.

Breadth gradient

The single test that most sharply distinguishes the two accounts is the breadth-stratified analysis in Figure 4. Under the projection reading we would expect roughly similar per-biome alignment magnitudes for biome-specific and universal motifs, since both are equally available as cultural projections of biome-correlated style. Under the coupling reading, only the biome-specific motifs are stable points of encounter with their biome’s visual structure; universals propagated precisely because they detached from any specific biome and stabilised on cross-cultural structures (family conflict, lost-then-found, transformation). The breadth-split result moves cleanly in the predicted direction. Marginal $\mu\Delta$ decays from $+0.52 \times 10^{-3}$ on Spec A motifs through $+0.26 \times 10^{-3}$ on semi-universals to $+0.05 \times 10^{-3}$ on universals; stratified $\mu\Delta$ decays from $+0.58 \times 10^{-3}$ through $+0.44 \times 10^{-3}$ to $+0.12 \times 10^{-3}$. Five of fourteen Spec A biome cells are individually significant under within-taxon stratification, against seven of fourteen for semi-universals and two of fourteen for universals; the Spec A subset carries a much larger per-biome effect




Figure 4. Breadth-stratified analysis. *Panel A:* across-biomes mean Δ on iNaturalist, both marginal and within-ionic-taxon stratified, partitioned by motif breadth into three groups (Spec A: motifs touching ≤ 3 biomes; semi-universal: 4–7 biomes; universal: ≥ 8 biomes). Both definitions show monotonic decay: stratified $\mu\Delta$ drops from $+0.58 \times 10^{-3}$ (Spec A) through $+0.44 \times 10^{-3}$ (semi-universal) to $+0.12 \times 10^{-3}$ (universal), the pattern the coupling account predicts and the projection account does not. *Panel B:* per-biome stratified Δ within each breadth group, with biomes ordered by mean stratified Δ across the three groups. Spec A bars (green) systematically dominate at the high- Δ biomes (Tropical Moist Broadleaf, Boreal/Taiga, Tropical Grasslands, Tundra, Tropical Dry Broadleaf); universal bars (red) systematically lie at or below zero across most biomes.

with proportionally similar significance density. The signature is present where the coupling account locates it, attenuated where breadth dilutes the environmental tie, and essentially absent where motifs have propagated across ≥ 8 biomes.

Per-taxon decomposition

Reading the per-biome $\times$ per-taxon matrix (Figure 5, Panel A) row by row reveals that every biome’s strongest cells fall in the taxa ecologically characteristic of it. Boreal Forests/Taiga is carried by Mammalia, Aves, and Actinopterygii (bears, owls, salmon); Tundra by Mammalia and Aves; Montane Grasslands by Mammalia, Reptilia, and Insecta; the wet tropics (Tropical Moist Broadleaf, Tropical Grasslands & Savannas) by Reptilia, Amphibia, Plantae, and Fungi; Mediterranean Forests almost exclusively by Plantae; Deserts & Xeric Shrublands by Mammalia, Plantae, and Reptilia. The alignment is not concentrated in any single class but distributes across the per-biome ecologically appropriate fauna and flora.

A trivial reading of this matrix is that the class words themselves are doing the work. Biome b’s photos contain biome b’s characteristic fauna, biome b’s anonymised motif text references those fauna by class word (“mammal” more often in boreal mythology, “reptile” more often in tropical mythology) because biome b’s mythology grew up around those animals, SigLIP notices the alignment as a class-word-frequency artefact, and the cells are positive by simple co-occurrence. The class-word collapse




Figure 5. Per-taxon decomposition with class-word collapse control. Biome $\times$ iconic-taxon residualised Δ on sentence-pooled SigLIP-2 over the full LLM-clean corpus. *Panel A*: class words preserved – every specific species name has been replaced by its class word (mammal / bird / fish / tree / plant, etc.) but the class words remain distinct. *Panel B*: class words collapsed – every animal-kingdom class word further reduced to the single token *animal* and every plant-kingdom word to *plant*, so class-word frequency is identical across biomes by construction. Rows are the 14 WWF biomes, ordered by descending pooled (all) Δ in Panel A. Cell values are per-biome per-taxon Δ in $\times 10^{-3}$, with stars for nominal significance and bold for BH- $q < 0.05$ within the taxon. Cells with fewer than 10 biome–taxon photographs are shown as grey em-dashes. The gold-bordered all column anchors the marginal headline. *Panel C*: for each taxon, a double-violin contrasts the distribution of its 14 biome cells under the two encodings – faded behind for Panel A, solid in front for Panel B. Taxa are ordered left-to-right by descending collapsed mean μ_C . Mean markers and per-taxon Wilcoxon p-values (against $\Delta = 0$) appear beneath. Plantae, Fungi, Aves, and Insecta retain most of their across-biome alignment under collapse, indicating that the per-(biome, taxon) coupling for those taxa is carried by surrounding lexical-thematic content (activities, landform vocabulary, narrative shape) rather than by the class word alone. Mammalia, Reptilia, and Amphibia attenuate substantially, indicating that a large part of their preserved-encoding signal was class-word frequency mirroring image-class frequency.

control in Panel B tests this directly. We re-embedded the same LLM-clean corpus after collapsing every animal-kingdom class word (mammal, bird, fish, reptile, amphibian, insect, mollusc, crustacean, arachnid) to the single token “animal” and every plant-kingdom word (tree, flower, plant, fungus) to “plant”. Class-word frequency is now constant across biomes by construction. Every animal mention reads “animal”, every plant mention reads “plant”, no biome’s text carries any class-word information beyond “some animal” or “some plant”.

The result is informative on both sides. The pooled alignment attenuates by about a third (the all column $\mu\Delta$ drops from $+0.244 \times 10^{-3}$ to $+0.160 \times 10^{-3}$, with the matrix retaining a per-cell Pearson $r = +0.57$ against its class-word-preserved counterpart). Per-taxon means change substantially for the vertebrate animal classes: Mammalia falls from $+0.82$ to $+0.09 \times 10^{-3}$, Reptilia from $+0.74$ to -0.16 , Amphibia from $+0.33$ to -0.14 . Plant content is more robust: Plantae drops only from $+0.61$ to $+0.48 \times 10^{-3}$. A substantial component of the per-(biome, taxon) signal for animals was class-word frequency mirroring image-class frequency. But the pooled effect does not collapse to zero, and the biomes whose cells stay strongly positive under the collapse (Tropical Moist Broadleaf, Temperate Broadleaf, Mediterranean Forests on Plantae and Fungi; Boreal Forests/Taiga and Tropical Coniferous Forests on Insecta and Mollusca) tell us where the coupling is carried by content beyond the class word.

The honest reading of the per-taxon matrix is therefore mixed. For animal-coded biomes, class-word frequency is doing a substantial fraction of the per-cell work, and the strong Mammalia and Reptilia cells

in Panel A should not be over- interpreted as evidence of biome-specific visual character of mammals or reptiles aligning with the text. For plant-coded and fungus-coded biomes, the alignment survives the collapse much better, and the surviving structure is more confidently attributable to the lexical-thematic content of the motif text aligning with how biome *b*'s actual flora appear in their habitat. The Kruskal–Wallis omnibus across taxa under either encoding is not significant, and no individual taxon survives BH correction at the taxon level; with only 14 biome cells per taxon, the test has limited power to discriminate distributions across taxa and we do not claim a statistical separation between them.

Discussion

The mythologies of nearly a thousand cultures, stripped of every species name, place name, biome word, and ethnonym we could identify, still align preferentially with images of the biomes those cultures inhabit. The alignment holds on two independent visual corpora, replicates across four vision–language models, and decays monotonically as motifs propagate beyond their biome of origin: tightest where the mythology is biome-specific, fading to nearly nothing where it has spread to most of the world. We read this pattern as evidence that mythology carries, in the lexical-thematic structure of its surviving text, a recoverable residue of the perceptual encounter between culture and biome.

Mythology as a recording layer of perceptual coupling

The breadth gradient is the result that most sharply distinguishes the two accounts of mythology we set out to test. A projection reading, on which mythology's symbolic content is largely a function of mind and culture and the environment supplies raw material but not form, has no principled reason to expect that biome-specific motifs and universal motifs should align with their biomes' imagery at different rates. The two classes of motif are equally available as projection surfaces. A coupling reading predicts the gradient directly: biome-specific motifs are the stable points of an encounter between a culture and a particular ecology, and their lexical-thematic content was shaped in that encounter; universal motifs propagated precisely because they detached from any single biome and stabilised on cross-cultural structures (family conflict, transformation, the journey to and from the dead) whose conditions for stabilisation are not ecological. Universals fade because the coupling has been factored out of them.

That reading converts mythology, in the technical sense, into a *recording layer*: not a description of the biome, and not a projection onto it, but a residue of the work the biome did in shaping which narrative completions were ecologically reinforced and which were not. The motif catalogues of computational mythography [Berezkin, 2024, ?, Tehrani, 2013] are, on this view, partial transcripts of a process that runs across many human generations and is irreducible to either side of the perceiver–environment pair. The phylogenetic mythology program [Tehrani, 2013, ?] has shown that motifs can be tracked across cultural genealogies as faithfully as genes; the breadth gradient we observe adds a complementary claim, that within those genealogies the motifs that stay tied to a single biome are the ones whose ecological substrate kept selecting for the same coupling, while the motifs that spread are the ones whose stabilisation was no longer biome-dependent.

A cognitive ecology of imagination

The lineage that makes this reading philosophically articulate runs across two centuries. Schelling's distinction between *natura naturata* (nature as finished product) and *natura naturans* (nature as productive activity becoming aware of itself) [Schelling, 1797, 1799] was recovered in the twentieth century by Bateson [Bateson, 1972, 1979], who turned it into a cybernetic claim: mind is wherever the conditions for a mental process obtain, wherever differences make differences, wherever pattern propagates across coupled systems. A human imagining a figure in a cliff face is not generating an inner image and applying it to outer matter. The imagining is an event in a coupled system that includes neural pattern-completion machinery, the actual statistical structure of the cliff, the cultural priors that have shaped what counts as

a figure to be found, and the linguistic apparatus that will or will not stabilise the percept into a transmissible name. Pull any element from the system and the event does not occur, or occurs differently. Depth psychology developed a parallel challenge to projection: Jung’s mature view located the archetypes in the joint between psyche and nature rather than in either alone [Jung, 1959]; Hillman read imagination as responsive, finding figures already there to be found [Hillman, 1975, 1981]; Corbin named the *mundus imaginalis*, a domain of real forms encountered through trained imagination and distinct from the modern collapse of “imaginary” into “merely subjective” [Corbin, 1972, Abram, 1996]. Ingold’s ecological anthropology generalised the move beyond personhood: human life is *dwelling*, not the projection of cultural schemes onto an external world but ongoing participation in a field of relations that produces both the dweller and the dwelt-in [Ingold, 2000, ?].

The cognitive sciences, beginning with Gibson’s ecological psychology [Gibson, 1979] and generalised by the 4E program (embodied, embedded, extended, and enactive cognition) [Varela et al., 1991, Clark and Chalmers, 1998, Noë, 2004], supplied the empirical scaffolding. Predictive processing gave the mechanism: the brain as a hierarchical prediction machine locking onto interpretations its priors privilege, given the actual statistical structure of what the world presents [Friston, 2010, Clark, 2016, Hohwy, 2013]. Niche construction theory closed the loop [Odling-Smee et al., 2003]: humans produce symbolic environments (names, stories, transmitted schemas) and those environments in turn shape the cognition of subsequent generations. On this synthesis, a figure in a biome is the stable point of a three-way coupling between environmental structure, perceptual machinery, and cultural transmission; imagination is not a private generative act but the moment of stabilisation in that coupling; and mythology is the cultural recording layer that keeps the stabilised completions available across generations.

The result reported here is the kind of evidence that lineage predicts and the projection model does not. A vision–language model is a system without perceptual phenomenology, without cultural participation, and without priors shaped by deep evolutionary time. That such a system, trained on the public products of the coupling, recovers the coupling’s signature in the joint statistical structure of environments and the stories told within them is what one would expect if the structure is really there to recover. It is, in Bateson’s vocabulary, a difference that makes a difference at the joint, detectable from either side.

What kind of coupling: the lexical-thematic level

A coupling claim has to be carried at some level of the text. Our analyses locate it at the lexical-thematic level rather than at the level of narrative structure. The alignment survives word-shuffling of each motif, so sentence-level composition is not where the residue lives; it disappears on synthetic encyclopedic paragraphs, so the residualised statistic does not manufacture a positive on biome-neutral content; it persists in motifs whose anonymised text alone provides minimal text-only biome predictability, so it is not reducible to a one-sided lexical artefact. What is doing the work, in other words, is the distribution of generic landform vocabulary and biome-correlated activity schemas that the noun-targeted anonymisation pipeline retained by construction.

This level is the one a strict perceptual-coupling reading predicts. Activities (gathering, hunting, fishing, herding, fording, climbing) are not free parameters of culture overlaid on an inert environment; they are biome-constrained possibilities, the affordances Gibson named, that the environment makes available to the perceiver. Generic landform words (forest, mountain, river, snow, sand, sea) are not biome labels but the articulated joints of the perceived world, the structural features around which attention and narrative organise [Ingold, 2000, ?]. A mythology that stabilises around the affordances and articulations of its biome is not failing to be symbolic; it is recording, in the only medium it has, the schemas of attention that the biome made available. The breadth gradient predicts that universals should be precisely the motifs whose stabilisation required affordances common to most environments, and that is what we find.

This is also the move that rescues the result from being merely methodological. The two-corpus replication is informative here. Activity vocabulary in the cleaned text could in principle align with activity *photographs* on a corpus that contained them. Places365’s strict-landscape filter removes humans, vehicles, and depictions of activity, leaving only terrain and vegetation, and the alignment persists with

comparable magnitude on that corpus. That is the signature of a coupling whose residue is read at the morphological level of the scene, not at the level of matching activity descriptions to activity images. The per-taxon decomposition makes a more careful version of the parallel argument from the species side. Each biome’s mythology aligns most strongly with the iconic taxa ecologically characteristic of that biome (Mammalia and Aves in the boreal and tundra, Reptilia and Amphibia in the wet tropics, Plantae in Mediterranean and temperate broadleaf forests). Under the maximally aggressive class-word collapse in which every animal-kingdom class word is reduced to “animal” and every plant-kingdom word to “plant”, the pooled alignment attenuates by about a third and the per-cell pattern correlates $r = +0.57$ with its uncollapsed counterpart, with substantial attenuation on the vertebrate animal classes and much smaller attenuation on the plant-coded and fungus-coded cells. A substantial component of the per-(biome, taxon) signal for animals was class-word frequency mirroring image-class frequency, and the strong Mammalia and Reptilia cells should not be read as direct evidence of biome-specific visual character of those animals aligning with the text. But the pooled effect does not collapse to zero, and the cells that survive the collapse (Plantae and Fungi across the moist tropics and temperate broadleaf forests; Insecta and Mollusca in the boreal and tropical coniferous belt) trace a coupling whose residue is carried by the surrounding lexical-thematic content of the anonymised motif, not by the class word alone. What gets recorded in the text and recovered by the model is partly the names of classes the biome made available and partly the perceptual shape of the biome that the mythology stabilised around.

Toward a falsifiable enactivist mythography

A philosophical reading that does not generate predictions is a worldview, not a research programme. The coupling lineage has historically been read as a worldview; the result reported here makes it possible to read it as a research programme with risky predictions. If imagination is coupling, and mythology is its recording layer, then several things should be true that the projection account either does not predict or actively predicts against. They should be increasingly directly testable.

First, the same coupling signature should appear in mythological *visual art* more strongly than in mythological text, because art is a more direct transcript of the perceptual encounter and has not been compressed through the narrative-summary bottleneck Berezkin’s catalogue records. A museum corpus is under construction. Second, the signature should be denser in oral-narrative audio than in literary transcription, because audio preserves more of the contextual texture that text drops. Third, the breadth gradient should mirror the historical distribution of subsistence ecology: motifs that decoupled from biome should track the cultural transmission routes (trade, migration, religious conversion) that moved them, with their original ecological substrate fading as they entered other biomes. Fourth, motif-conditioned visual priors should improve scene classification on photographs of the biomes those motifs were stabilised in, because the priors encode information about which configurations were perceptually salient in that biome. None of these is a tautology; each can fail; and if mythology is doing what we suggest it is doing, each should pass.

The honest reading of what the data show is bounded. They do not demonstrate that nature is conscious, or that mythologies are messages from a non-human agency, or that the strongest metaphysical readings of the coupling lineage (Whitehead’s process metaphysics [Whitehead, 1929], Goff’s panpsychism or Strawson’s realistic monism [Goff, 2019, Strawson, 2006]) are empirically supported. They show that the joint statistical structure of environments and the stories told within them carries a recoverable residue of their coupled history, robust enough that a system without phenomenology can detect it from either side. That is a smaller claim than the strongest readings of the lineage and a larger claim than the projection account allows.

Berezkin’s corpus is curated, Eurasian-weighted, and sparse in several biomes (Mangroves, Flooded Grasslands, Temperate Conifer Forests). The vision–language models we use carry the biases of their training distributions, which the cross-model agreement bounds but does not erase. The anonymisation pipeline is noun-targeted; an activity-vocabulary pass that strips the remaining biome-correlated verbs is the natural next experiment, and we expect it to attenuate but not eliminate the signal, on the same logic that the within-Glottolog biome-swap null attenuated rather than eliminated it: the biome-content

information is not concentrated in any single layer of the text but distributed across the substrate that the perceptual encounter shaped.

Schelling wrote that “nature should be visible Mind, and Mind invisible Nature.” The pattern reported here is not a demonstration of his stronger thesis, but it is the kind of evidence that thesis predicts and the modern projection account does not. Mythology, on the reading the data support, is one of the long instruments humans have used to keep the coupling articulate. The work of making the coupling reading testable has only begun.

Materials and Methods

Corpus and biome assignment

We use Yu. E. Berezkin’s *Analytical Catalogue of World Mythology and Folklore* [Berezkin, 2024], scraped from `ruthenia.ru` in May 2024. The catalogue provides 2,564 motif entries indexed across 958 traditions; each motif has a multi-paragraph Russian-language abstract aggregating the variants found across the traditions that carry the motif. Tradition centroids are joined to WWF Terrestrial Ecoregions [Olson et al., 2001] to assign each tradition, and through it each tradition–motif pair, to one of the 14 WWF biomes.

Image corpora

The iNaturalist component consists of 47,478 images drawn from research-grade observations on iNaturalist Open Data [iNaturalist, 2024] across the 14 WWF biomes. We apply an image-side sanity filter that drops obvious non-nature content via SigLIP-2 zero-shot scoring against twelve negative prompts (“a person in the photo”, “a portrait of a human”, “people holding something”, “an angler holding fish”, “an indoor scene”, “a kitchen or laboratory”, “a city skyline or street”, “a car or vehicle”, “a road sign”, “a black or all-white image”, “a museum specimen on a white background”, “a screenshot or document”). An image is dropped if its maximum cosine similarity to any of the twelve negative prompts exceeds 0.08. The threshold is lenient (removing only obvious cases) so that species-in-habitat close-ups and natural outdoor scenes are retained, which is the content iNaturalist is biased toward by design. For the curated landscape comparison, the Places365 component [Zhou et al., 2018] is drawn from scene categories whose semantic interpretation unambiguously corresponds to a single WWF biome (e.g. *tundra* → arctic tundra; *rainforest* → tropical moist broadleaf; *forest/broadleaf* → temperate or tropical dry broadleaf; *snowfield/mountain_snowy* → boreal/taiga; *lagoon/swamp/creek/coast* → mangroves); categories that mix biomes were excluded. Each candidate image is then passed through a strict-landscape SigLIP-2 filter that drops anything that is not a wide-shot natural scene: five positive prompts (“a wide-shot natural landscape with no people”, “a panoramic photograph of natural scenery”, “an empty natural habitat photographed from a distance”, “a landscape photo of natural terrain showing vegetation and horizon”, “a wild ecosystem with no human-made structures”) and eleven negative prompts (“a person in the photograph”, “a portrait of one or more humans”, “a close-up photo of a single animal”, “a single bird filling the frame”, “a car or vehicle”, “a boat”, “a road or street”, “a man-made structure or building”, “an indoor scene”, “a person riding a horse or animal”, “a tourist photograph with people”) are encoded by SigLIP-2 and compared against each image’s embedding. The strict-landscape score is $\text{mean}(\text{positives}) - \text{max}(\text{negatives})$; an image is dropped if its $\text{max}(\text{negatives}) > 0.18$, and the remaining top- K images per biome are kept. The Places365 filter is intentionally stricter than the iNaturalist filter, in particular it removes single-species close-ups, because the Places365 panel functions as a scene-only robustness check on biome-character imagery, complementary to iNaturalist’s species-in-habitat coverage.

Vision–language embedding

The headline model is SigLIP-2-large (google/siglip2-large-patch16-384) [Tschannen et al., 2025], a multilingual vision–language model with native Russian support. Images use the standard 384 × 384 preprocessing. Text is sentence-pooled: each LLM-clean Berezkin abstract is split into sentences (regex on terminal punctuation, sentences of < 3 words dropped, capped at 50 sentences per motif), each sentence is embedded independently within SigLIP-2’s 64-token window (mean 23 sentences per motif; 50,032 unique sentences across 2,158 motifs in total), and the per-motif vector is the mean-pooled, L2-renormalised sentence embedding. Sentence pooling matters because the raw English abstracts average ~ 3,400 tokens, of which a 64-token truncation captures only the first ~ 2%; the sentence-pooling pass embeds the full anonymised abstract instead. For cross-model robustness we additionally embed with M-CLIP (XLM-Roberta-Large + CLIP-ViT-L/14) [Carlsson et al., 2022] on the Russian abstracts (M-CLIP’s 512-token window covers most abstracts in a single forward pass) and with OpenCLIP-ViT-L/14 (LAION-2B and OpenAI pretrainings) on English motif oneliners (see Supplementary Figure 9). Image and text features are L2-normalised and similarity is the inner product.

Controls against biome cues




Figure 6. Anonymisation pipeline and image-side stratification against biome cues. *Stage 1.* Gemini 3.5 Flash translates each Russian Berezkin abstract into English under seven structured rules (species → class word; place → [place]; ethnonym → [people]; biome words dropped; crops → “plant”; cultural artefacts → generic; Latin binomials dropped). *Stage 2.* A second Gemini pass standardises long-tail cultural titles (sultan, kadi, padishah → “ruler” / “holy person”) and tradition-specific character names (Baba Yaga, Kivioq, Coyote-character → “character”) to neutral generics; held-out audit confirms a 80–100% clean rate. *Stage 3.* Each cleaned text is embedded by SigLIP-2 and scored against the biome visual prototypes (mean image embedding per biome). Motifs whose tell $z(M)$ exceeds the 95th percentile of a 1,000-permutation scrambled-label null are dropped ($z_{95} = 2.83$, 237/2,158 motifs excluded). *Stage 4.* The Δ test is recomputed within each iconic-taxon stratum on the image side and averaged uniformly, removing biome by taxon co-occurrence as a possible explanation of alignment. The displayed example is motif k87b (“Woman steps into bears’ dung”, Boreal/Taiga).

The motif text fed to the encoder is anonymised against every category of biome-anchoring vocabulary we could identify, by an LLM-mediated translation and refinement pass followed by a single exclusion filter. For each motif we send the raw Russian Berezkin abstract together with the original English oneliner to gemini-3.5-flash (Google’s frontier flash model, released May 2026) [?] with a structured prompt that asks it to produce a faithful English translation under the following rules:

- Every specific species name → its class word: one of mammal / bird / fish / reptile / amphibian / insect / mollusc / crustacean / arachnid / tree / flower / plant / fungus / animal.
- Every country, region, river, mountain, lake, sea, ocean, or city name → the literal string [place].

- Every tribe, ethnic group, nation, or ethno-religious community name → [people].
- Every biome word (tundra, taiga, savanna, jungle, rainforest, prairie, desert, steppe, swamp, marsh, oasis, glacier, veldt, mangrove, mediterranean, etc.) dropped.
- Every cultivated crop (rice, corn, beans, peanuts, millet, quinoa, cassava, cotton, indigo, tobacco, etc.) → plant.
- Every biome-coded cultural artefact (sledge, kayak, harpoon, igloo, yurt, blowgun, snowshoes, kachina, calabash) → a generic equivalent (tool, boat, shelter, vessel).
- Latin binomials dropped.

A second Gemini pass applied to the first-pass output catches the long tail and additionally:

- Cultural-religious titles (sultan / padishah / shah / kadi / mullah / dervish / brahmin / bai / lama / raja / khan / tsar) → neutral “ruler” or “holy person” or “wealthy person”, explicitly *not* the Western-coded “king” / “queen” / “priest” / “monk”.
- Tradition-specific named characters (Baba Yaga, Kivioq, Coyote-as-character, Nach Ak Yum, Quetzalcoatl, Khevki, Mitoshka, etc.) → “character”.

Kept by design as cross-cultural anchors: cross-cultural religious figures (Christ, Buddha, Adam, Eve, Allah, God, Satan, named saints, Krishna, Rama, Sita), universal personifications (Sun, Moon, Wind, Sky, Earth, Sea, Thunder, Lightning, Star), all astronomical references (Pleiades, Orion, Big Dipper, Sirius, Venus, Milky Way, named constellations), all class words, cardinal directions, and generic universal vocabulary (forest, mountain, river, village, house, person, woman, man, child, day, night, fire, water). The narrative content of each motif, transformations, family relationships, magical events, plot, is preserved.

A held-out audit of 200 random LLM-anonymised motifs (criteria itemised in the Results panel C) finds 0% explicit biome words, 0.5% place or ethnonym residues, and 3.0% specific species residues, giving a composite strict-clean rate of 96.5% against species, place, and biome vocabulary. Generic landform vocabulary (forest, mountain, river, snow, sand, 92.5% of motifs) and biome-correlated activity verbs (gather, hunt, fish, herd, 93.5% of motifs) are retained by construction: the anonymisation pipeline is *noun-targeted*, stripping species, place, biome, and cultural-artefact nouns while leaving narrative structure and activity predicates intact. A residual narrative such as “gathers fruit, steps in mammal dung” therefore retains activity-level biome correlates (gathering fruit is more frequent in temperate and boreal forest descriptions than in desert ones). We treat this retained content as the substrate over which the alignment is measured rather than as contamination, and return to the morphology-versus-activity-priming question in the Discussion.

Within-taxon image-side stratification.

Because iNaturalist observations are taxonomically structured, some biomes are disproportionately rich in particular iconic taxa (boreal/tundra/conifer forests are 55–66% Plantae; mangroves only 11%). A motif whose anonymised text emphasises one taxon class will therefore score high in any biome whose iNat image pool emphasises the same taxon class, even when the biome’s specific visual character contributes nothing. To remove this biome-by-taxon co-occurrence as a possible explanation for the alignment, we recompute Δ_b stratified by the iNaturalist iconic-taxon label. For each biome b and taxon stratum t with at least 20 images, we compute the within-stratum $\Delta_{b,t} = \overline{S_{i, \mathcal{M}_{bi} \in b \cap t}} - \overline{S_{i, \mathcal{M}_{bi} \in b \cap t^c}}$, and report the uniform across-strata mean as $\Delta_b^{\text{strat}} = (1/T_b) \sum_t \Delta_{b,t}$. The permutation null is the same scrambled-motif null applied within strata. The headline iNaturalist statistic reported throughout the main text and main figures is stratified ($\mu\Delta^{\text{strat}} = +0.406 \times 10^{-3}$ across the 14 biomes), with the marginal $\mu\Delta^{\text{marg}} = +0.244 \times 10^{-3}$ reported alongside in robustness panels (where the input rather than the statistic is being varied) and on the Places365 panel of Figure 2 (Places365 has no iconic-taxon labels and therefore admits no further stratification). Per-text-taxon decompositions (Figure 5) report the per-cell marginal value at each

biome–taxon cell, since each cell already conditions on a single taxon; the stratified Δ is, mechanically, the uniform mean of those per-cell marginals.

We separately explored simpler rule-based anonymisation pipelines (genus-level hypernymisation; a hand-curated class-level Russian hypernym map of roughly 3,500 entries; an audit-extended map of 6,300 entries) and found that the effect is positive under each of them and the across-biome ranking is preserved, ruling out the interpretation that any particular vocabulary anchor is doing the work. The LLM-mediated pipeline described above is the strictest in the family and the one we use for the headline.

Statistical tests

Image–motif similarities form an $n_{\text{img}} \times n_{\text{motif}}$ matrix $\mathbf{S}$. Per-motif residualisation (subtracting the motif grand mean) eliminates the contribution of generic motif valence to per-biome aggregates. For each biome b with n_b images and motif partition $(\mathcal{M}_b, \mathcal{M}_{\bar{b}})$ where $\mathcal{M}_b$ is the set of motifs whose Berezkin tradition set includes at least one tradition assigned to biome b , we report

$$\Delta_b = \frac{1}{n_b} \sum_{i \in \text{imgs}(b)} \left[\overline{S_{i, \mathcal{M}_b}} - \overline{S_{i, \mathcal{M}_{\bar{b}}}} \right],$$

the residualised marginal Δ at the biome level. The headline iNat statistic is the within-iconic-taxon stratified version Δ_b^{strat} defined above. Significance is assessed by 1,000 permutations of the motif-to-biome assignment (with the stratified statistic recomputed under each permutation), with FDR control across biomes via the Benjamini–Hochberg procedure [Benjamini and Hochberg, 1995]. For all main-figure panels we apply a display threshold of ≥ 50 in-biome images, ≥ 5 in-biome motifs, and ≥ 20 images per biome–taxon stratum; biome–taxon cells with fewer than 10 images appear as missing in the per-taxon decomposition (Figure 5).

Data and code availability

All code, sentence-pooled SigLIP-2 embeddings, the LLM-clean motif texts (Gemini-3.5-Flash output), the held-out anonymisation audit, and the earlier rule-based hypernym maps released for completeness are available at [repository URL placeholder]. The Berezkin abstracts are accessible at <http://www.ruthenia.ru/folklore/berezkin/>.

Acknowledgements

[Placeholders.]

References

- David Abram. *The Spell of the Sensuous: Perception and Language in a More-Than-Human World*. Pantheon Books, 1996.
- Owen Barfield. *Saving the Appearances: A Study in Idolatry*. Faber and Faber, 1957.
- Gregory Bateson. *Steps to an Ecology of Mind*. Ballantine Books, 1972.
- Gregory Bateson. *Mind and Nature: A Necessary Unity*. Dutton, 1979.
- Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society B*, 57(1):289–300, 1995.
- Yuri E. Berezkin. Analytical catalogue of world mythology and folklore. <http://www.ruthenia.ru/folklore/berezkin/>, 2024.

- Fredrik Carlsson et al. Cross-lingual and multilingual CLIP. arXiv preprint arXiv:2207.07551, 2022.
- Andy Clark. *Surfing Uncertainty: Prediction, Action, and the Embodied Mind*. Oxford University Press, 2016.
- Andy Clark and David Chalmers. The extended mind. *Analysis*, 58(1):7–19, 1998.
- Henry Corbin. Mundus imaginalis, or the imaginary and the imaginal. *Spring*, pages 1–19, 1972.
- Karl Friston. The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11: 127–138, 2010.
- James J. Gibson. *The Ecological Approach to Visual Perception*. Houghton Mifflin, 1979.
- Philip Goff. *Galileo’s Error: Foundations for a New Science of Consciousness*. Pantheon Books, 2019.
- James Hillman. *Re-Visioning Psychology*. Harper & Row, 1975.
- James Hillman. *The Thought of the Heart and the Soul of the World*. Spring Publications, 1981.
- Jakob Hohwy. *The Predictive Mind*. Oxford University Press, 2013.
- iNaturalist. iNaturalist Open Data on AWS. <https://github.com/inaturalist/inaturalist-open-data>, 2024.
- Tim Ingold. *The Perception of the Environment*. Routledge, 2000.
- Carl Gustav Jung. *The Archetypes and the Collective Unconscious*. Collected Works of C. G. Jung, Vol. 9, Part 1. Princeton University Press, 1959.
- Alva Noë. *Action in Perception*. MIT Press, 2004.
- F. John Odling-Smee, Kevin N. Laland, and Marcus W. Feldman. *Niche Construction: The Neglected Process in Evolution*. Princeton University Press, 2003.
- David M. Olson et al. Terrestrial ecoregions of the world: A new map of life on earth. *BioScience*, 51 (11):933–938, 2001.
- Friedrich Wilhelm Joseph Schelling. *Ideas for a Philosophy of Nature*. Cambridge University Press, 1797. English ed. 1988, trans. Errol E. Harris and Peter Heath.
- Friedrich Wilhelm Joseph Schelling. *First Outline of a System of the Philosophy of Nature*. State University of New York Press, 1799. English ed. 2004, trans. Keith R. Peterson.
- Galen Strawson. Realistic monism: Why physicalism entails panpsychism. *Journal of Consciousness Studies*, 13(10–11):3–31, 2006.
- Jamshid J. Tehrani. The phylogeny of little red riding hood. *PLoS ONE*, 8(11):e78871, 2013.
- Michael Tschannen et al. Siglip 2: Multilingual vision–language encoders with improved semantic understanding, localization, and dense features. arXiv preprint arXiv:2502.14786, 2025.
- Francisco J. Varela, Evan Thompson, and Eleanor Rosch. *The Embodied Mind: Cognitive Science and Human Experience*. MIT Press, 1991.
- Alfred North Whitehead. *Process and Reality*. Free Press, 1929. Corrected ed. 1978.
- Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba. Places: A 10 million image database for scene recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(6):1452–1464, 2018.

Supplementary materials

The supplementary materials below collect the analyses that probe the robustness of the headline finding from five complementary angles: (i) the headline robustness atlas summarised in the main Results, with the four panels shown at full size; (ii) a per-taxon decomposition that asks which iNaturalist taxonomic groups carry the alignment; (iii) a cross-model replication across four vision–language model architectures on the same LLM-clean text; (iv) two analyses directly addressing whether biome-correlated lexical content in the cleaned text could by itself explain the alignment (a high-tell versus low-tell split and a within-Glottolog-macroarea biome-swap null); and (v) an analysis of the raw ecological composition of motif text per biome, which makes explicit what the anonymisation pipeline is removing. Throughout, the same statistical conventions apply as in the main text: sentence-pooled SigLIP-2 on the full LLM-clean corpus, residualised Δ with 1,000 motif-to-biome permutations and Benjamini–Hochberg FDR within each panel, and gold marks for $q < 0.05$.

S1. Robustness atlas

Four robustness analyses bound where the headline signal comes from. Word-shuffling each motif before sentence-pooling leaves the stratified $\mu\Delta$ essentially unchanged, indicating that the alignment depends on lexical content rather than narrative structure. Replacing the motifs with synthetic encyclopedic-register paragraphs of comparable length collapses $\mu\Delta$ to near zero, ruling out the possibility that the residualised statistic produces a positive bias on biome-neutral text. The per-biome Δ also replicates across all four vision–language model architectures we tested (SigLIP-2 sentence-pooled, M-CLIP, OpenCLIP-LAION-2B, OpenCLIP-OpenAI) with across-model Pearson correlations between +0.74 and +0.94, ruling out a SigLIP-specific training-distribution artefact. Finally, a held-out audit of 200 random anonymised motifs confirms a strict-clean rate of 96.5% against species, place, and biome vocabulary; the categories that remain by construction (generic landform words, activity verbs) are the substrate over which we measure alignment, and we interpret them as such in the main text Discussion.

S2. Per-taxon decomposition (small-multiples view)

Figure 8 re-renders the per-taxon decomposition of the headline result as small-multiple bar panels, one taxon per panel, with within-taxon BH-FDR. This is a complementary view of the same data as Figure 5: rather than collapsing a taxon’s twelve biome cells into one violin, each panel shows the per-biome Δ values themselves, so the reader can see which specific biomes carry each taxon’s effect. The “all” panel reproduces the headline ranking (Tropical Grasslands, Mediterranean, Boreal/Taiga at the top); the Aves and Plantae panels echo this pattern; the Reptilia, Mollusca, and Arachnida panels show little structure, consistent with the per-taxon violin summary.

S3. Cross-model replication

Figure 9 tests whether the per-biome stratified Δ pattern is an artefact of the specific vision–language model architecture used to produce the embeddings (sentence-pooled SigLIP-2-large in the headline analysis). For this robustness check we re-run the within-iconic- taxon stratified pipeline with four distinct vision–language model families on the same LLM-clean English motif text: sentence-pooled SigLIP-2-large; M-CLIP (XLM-Roberta-Large paired with CLIP-ViT-L/14, 512-token native window); OpenCLIP-ViT-L/14-LAION-2B; and OpenCLIP-ViT-L/14-OpenAI. The 4×4 scatter matrix shows per-biome stratified Δ correlations for every model pair; the off-diagonal cells report Pearson r . Correlations lie in $[+0.74, +0.94]$ across all six distinct model pairs, with the SigLIP-2-vs-OpenCLIP-LAION-2B pair carrying the strongest agreement ($r = +0.94$). The per-biome ranking is preserved across every model family tested, ruling out reading on which the headline pattern is a SigLIP-2-specific artefact of one model’s training distribution.



Figure 7. Robustness atlas referenced in the main Results. *Panel A:* per-biome stratified Δ on iNaturalist under three conditions that modify the text input. The original sentence-pooled embedding (dark blue) is essentially identical to its word-shuffled version (light blue) at every biome, indicating that the alignment depends on lexical content rather than narrative structure. The encyclopedic null (grey), in which synthetic encyclopedic-register English paragraphs of comparable length replace the LLM-clean motifs, collapses the alignment toward zero across biomes. *Panel B:* per-biome stratified Δ on iNaturalist species photographs (x-axis) versus marginal Δ on Places365 landscape scenes (y-axis), one point per WWF biome that both corpora cover. Pearson $r = +0.36$ across the 11 shared biomes; mean Δ positive on both sides ($+0.46 \times 10^{-3}$ iNat-shared stratified, $+0.42 \times 10^{-3}$ Places365 marginal). The two corpora share no images and no visual content type, so this is an independent cross-channel replication. *Panel C:* expanded 200-motif audit, per-criterion residue rate. Biome words, place or ethnonym residues, and specific species references are essentially eliminated (0%, 0.5%, and 3.0% respectively; composite strict-clean rate 96.5%). Generic landform vocabulary and biome-correlated activity verbs are retained by design. *Panel D:* cross-model headline stratified $\mu\Delta$ on the same LLM-clean motif text. Stratified $\mu\Delta$ is positive on all four architectures, from $+0.41 \times 10^{-3}$ on SigLIP-2 to $+0.74 \times 10^{-3}$ on OpenCLIP-LAION-2B.




Figure 8. Per-taxon biome panels. Per-taxon decomposition of the headline result, presented as small-multiple bar panels with within-taxon BH-FDR. Each panel shows the per-biome Δ for one taxon; biome row order is preserved across panels. Gold stars mark BH- $q < 0.05$ within the panel. Complements Figure 5.

S4. Biome-tell tests

A residual concern about the headline alignment is that the anonymised text might still contain enough biome-correlated vocabulary (generic landform words, activity verbs, class-level animal terms) for SigLIP-2 to recover biome assignment from text alone, with the image side being a secondary recovery of that lexical signal rather than an independent confirmation. We address this concern with two analyses whose details are provided in the data drop.

High-tell versus low-tell split. For each motif we compute a biome-tell z -score: the gap between its maximum sentence-pooled SigLIP-2 similarity to its assigned biomes’ visual prototypes and the mean similarity to non-assigned biomes’ prototypes, normalised by the standard deviation of off-target similarities. We split motifs at the median tell z -score and run the within-iconic-taxon stratified Δ test on each half (Table 1). The high-tell half (motifs whose anonymised text *is* biome-recoverable above chance) gives $\mu\Delta^{\text{strat}} = +0.54 \times 10^{-3}$ with 7 of 14 biomes significant at $p < .05$; the low-tell half (motifs whose anonymised text alone is *least* biome-recoverable) still gives $\mu\Delta^{\text{strat}} = +0.21 \times 10^{-3}$ with 7 of 14 biomes significant at $p < .05$. Text-only biome predictability therefore contributes to the alignment (the high-tell half is about $2.5\times$ the low-tell half), but the low-tell half on its own is not zero. That motifs whose text alone provides minimal biome predictability still align positively with their biome’s imagery indicates that the headline signal is not fully explained by biome-correlated lexical content recoverable




Figure 9. Cross-model robustness. 4×4 scatter matrix of per-biome within-iconic-taxon stratified Δ on the same LLM-clean Berezkin abstracts and the same iNaturalist images, across four vision–language model architectures. Off-diagonal cells report Pearson r between models; the dashed line marks $y = x$. Each point is one of the 14 WWF biomes, coloured by the standard biome palette. All six distinct model pairs correlate in $[+0.74, +0.94]$.

from text alone.

Table 1. High-tell versus low-tell split. Median-split of motifs on the biome-tell z -score (text-only biome predictability), with the within-iconic-taxon stratified Δ test recomputed on each half on iNaturalist images. Both halves carry a positive across-biome mean Δ .

Subset	$\mu\Delta^{\text{strat}} (\times 10^{-3})$	$\mu\Delta^{\text{marg}} (\times 10^{-3})$	Sig $p < .05$
Low-tell half	+0.21	−0.41	7/14
High-tell half	+0.54	+0.78	7/14

Within-Glottolog-macroarea biome-swap null. A stronger version of the macro-area block-permutation null is to actively swap biome labels while preserving cultural geography. For each tradition we ran-

domly reassign its motifs to another tradition in the same Glottolog macro-area but a *different* biome, and recompute the per-biome stratified Δ under each of 1,000 such swaps. If biome assignment is doing the work, the swapped Δ should collapse below the observed value; if biome identity is irrelevant and the alignment is purely a within-macro-area lexical artefact, the swap distribution should contain the observed value. Six biomes carry observed stratified Δ above the 95th percentile of the swap null distribution: Montane Grasslands & Shrublands (observed $+1.08 \times 10^{-3}$, swap-null mean $+0.09$, $p < .001$), Tropical & Subtropical Moist Broadleaf Forests (observed $+0.93$, null mean -0.16 , $p < .001$), Boreal Forests/Taiga (observed $+0.65$, null mean $+0.08$, $p < .001$), Deserts & Xeric Shrublands ($+0.63$ vs null $+0.22$, $p = .018$), Tropical & Subtropical Grasslands & Savannas ($+0.76$ vs null $+0.45$, $p = .019$), and Temperate Grasslands, Savannas & Shrublands ($+0.33$ vs null $+0.02$, $p = .022$). The remaining biomes either have observed values indistinguishable from the swap null (Tundra, Mediterranean, Temperate Conifer, Temperate Broadleaf) or are too sparsely represented for a stable swap distribution (Mangroves, Flooded Grasslands). This is a substantially stronger test than the standard motif-to-biome permutation, because it preserves within-macro-area cultural autocorrelation by construction: the biome label is being shuffled among traditions that share a linguistic-cultural region. That six biomes still carry signal above this stricter null indicates the alignment is not fully explained by within-macro-area lexical regularities — the biome content of the motif text itself carries information beyond what would be recoverable from cultural-region membership alone.




Figure 10. Biome-tell tests. *Panel A:* per-biome within-iconic-taxon stratified Δ on the high-tell half (motifs whose anonymised text is biome-recoverable above chance, gold bars) and the low-tell half (motifs whose text alone is biome-uninformative, blue bars), across the 14 WWF biomes. Both halves carry positive across-biome mean Δ ($+0.54$ and $+0.21 \times 10^{-3}$ respectively); seven biomes are individually significant at $p < .05$ in each half. *Panel B:* observed stratified Δ per biome (coloured markers) overlaid on the 95% range of the within-Glottolog-macroarea biome-swap null distribution (grey strips). Circles mark biomes whose observed Δ exceeds the 95th percentile of the swap null ($p < .05$); squares mark biomes that do not. Six biomes survive this stricter null, indicating that biome content carries information beyond within-macro-area cultural autocorrelation.

S5. Raw ecological composition per biome (Spec A view)

Figure 11 reports the raw ecological composition of motif text per biome, before anonymisation. For each WWF biome (rows) and each of fourteen taxonomic and botanical categories (columns), the left panel of the heatmap shows the percentage of Spec A motifs whose raw Russian abstract mentions at least one species or plant of that category. The right panel shows the difference between the Spec A subset and the all-in-biome pool: a positive cell means the category is over-represented in biome-specific motifs relative to the universal-dominated full pool. The heatmap therefore makes explicit, at the textual level, which taxonomic categories carry biome-specific signal before the hypernymisation layer removes their

names from the encoder’s input. The figure helps motivate which species-vocabulary categories the v7 hypernym map needed to cover comprehensively, and clarifies the ecological content the encoder is no longer reading at the headline stage.




Figure 11. Ecological composition per biome (Spec A subset). *Left:* for each WWF biome (rows) and each of 14 taxonomic and botanical categories (columns), the cell shows the percentage of *Spec A motifs* (motifs touching ≤ 3 biomes, ≥ 3 own-traditions in that biome, the same selection that drives the headline result) whose raw Russian abstract mentions at least one species or plant of that category, before anonymisation. *Right:* difference (Spec A minus all-in-biome). Positive cells mark categories over-represented in biome-specific motifs relative to the universal-dominated full pool, revealing the ecological fingerprint that the species-hypernymisation layer subsequently removes. The all-motifs version of the heatmap is dominated by universal motifs and therefore uniformly red across biomes; we omit it here and release the underlying counts in the data drop.